



Open Source Enterprise Integration Platform

*Jeff Mitchell
ikasan.org
July 2010*

Introduction



- The Problem Domain
- Architecture Strategies
- Translating the Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Problem Domain



- **[Domain]** Enterprise Application Integration (EAI) space
 - Integration of business systems in a landscape of interwoven and often complex finance processes
- **[Issue]** EAI is complex, costly, & not the primary business concern
 - Greater adoption of “Best of Breed” specialist applications
 - Business applications distributed across disparate platforms
 - Business data distributed across isolated silos
 - Legacy data repositories
 - Data duplication and integrity issues
 - No clear business data owner
 - Exponential integration requirements
 - Greater complexity of business demand on data orchestration
 - EAI is more than simply connecting applications
- **[Goal]** To provide simple, robust, configurable commoditised solutions
 - Expose business artifacts whilst isolating the integration specifics

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Architecture Strategies



- Simplicity
 - Must not require in-depth proprietary knowledge
 - Avoidance of overly complex or heavyweight frameworks
 - Highly testable simple constructs
- Commonality
 - Reusable interchangeable constructs
 - Standard contracts of interaction
 - Simple repeatable implementation steps
- Adaptability
 - Ability to support any type of business entity
 - Ability to integrate any type of Enterprise Information System
 - Without breaking the first two strategies of Simplicity & Commonality
- Robust & Guaranteed operation
 - Maintain data integrity in business delivery and failure scenarios
 - Require minimal manual intervention i.e. failure / automated recovery

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Architecture Strategies



- Clear contract definition and separation of concerns

- Problem Domain
- **Architecture Strategies**
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Architecture Strategies



- Clear contract definition and separation of concerns
 - Application concerns

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary



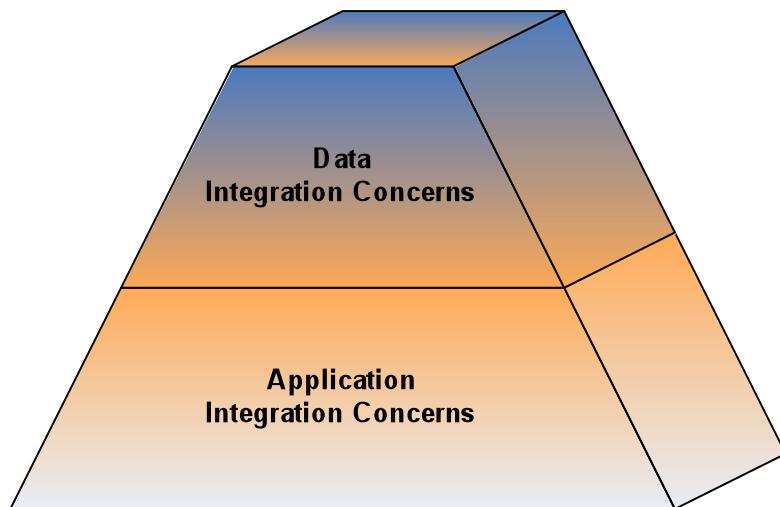
- **[Application]** performant, robust & guaranteed operation

Architecture Strategies



- Clear contract definition and separation of concerns
 - Application concerns
 - Data concerns

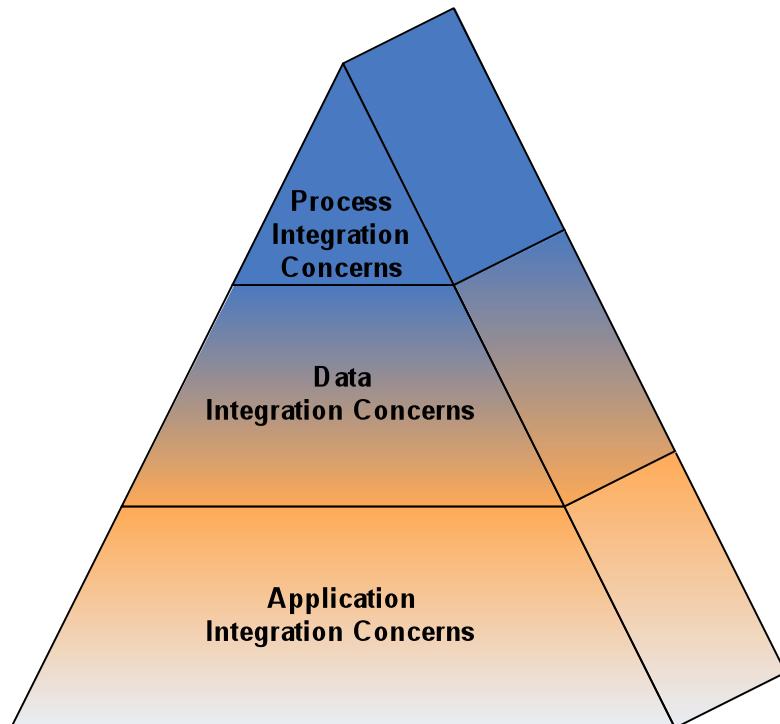
- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary



- **[Data]** presentation of standard, meaningful business data
- **[Application]** performant, robust & guaranteed operation

Architecture Strategies

- Clear contract definition and separation of concerns
 - Application concerns
 - Data concerns
 - Business process concerns



- Problem Domain
 - Architecture Strategies
 - Translating Architecture
 - Technology Stack
 - Features
 - Development Lifecycle
 - Roadmap
 - Summary
-
- **[Process]** business event & entity orchestration (STP is the goal)
 - **[Data]** presentation of standard, meaningful business data
 - **[Application]** performant, robust & guaranteed operation

Architecture Strategies



- Loose coupling
 - Integrations should only require knowledge of each other through exchange of business artifacts
- Tight cohesion
 - An integration's specifics must not bleed out unnecessarily
- High visibility
 - Business data tracking
 - Fail scenarios
 - Operational audit
- Single logical point of integration
 - Geographically agnostic / cross platform support
- Open standards aligned with proven design patterns
 - Adoption of industry standard Enterprise Design Patterns
 - Avoidance of vendor lock-in

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Architecture Strategies



[Key Strategies]

▪ Single Point of Integration	▪ Loose Coupling	▪ Variety of Standard Contracts
▪ Separation of Concerns	▪ Tight Cohesion	▪ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

- Ensure traceability from architecture through to implementation
 - All too often architecture fails to translate into the real world
 - The right architecture can reduce complexity

Translating the Architecture



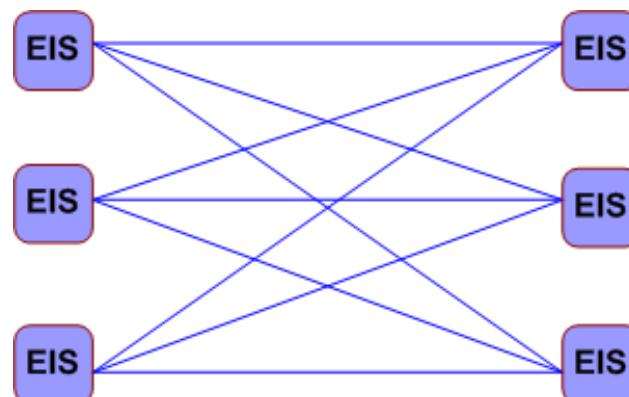
[Key Strategies]

▪ Single Point of Integration	▪ Loose Coupling	▪ Variety of Standard Contracts
▪ Separation of Concerns	▪ Tight Cohesion	▪ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Issue] Commonly agreed that “spaghetti integration” is bad

- Integration can be exponentially complex
- Bleed out of APIs, data syntax and business semantics
- Ripple effect of change is massive



Translating the Architecture



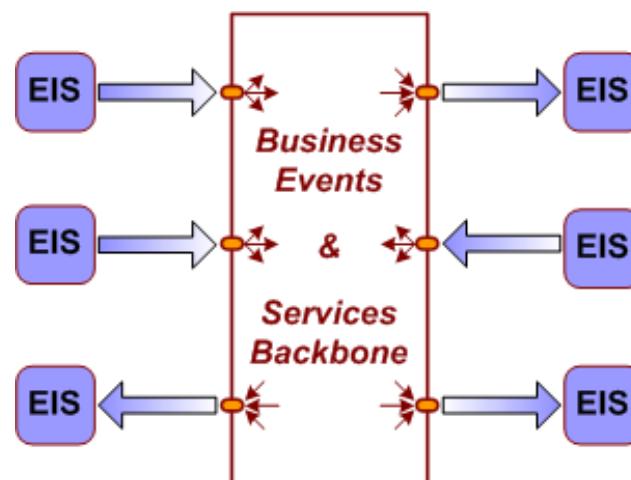
[Key Strategies]

- | | | |
|-------------------------------|------------------------------|---------------------------------|
| ▪ Single Point of Integration | ▪ Loose Coupling | ▪ Variety of Standard Contracts |
| ▪ Separation of Concerns | ▪ Tight Cohesion | ▪ Standard Design Patterns |
| ▪ Interchangeable Constructs | ▪ Highly Testable Constructs | ▪ Data Integrity |
| ▪ Control & Management | ▪ Monitor & Alerting | ▪ Minimal Manual Intervention |
| ▪ Support any Protocol/Entity | ▪ Business Event Tracking | ▪ Operational Audit |

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Concept] Standard Enterprise Application Integration (EAI) approach

- Provision of an Event/Service backbone
- Single point of integration for EIS business flow
- Ripple effect of change is localised (*assuming best practice*)



Translating the Architecture



[Key Strategies]

▪ Single Point of Integration	▪ Loose Coupling	▪ Variety of Standard Contracts
▪ Separation of Concerns	▪ Tight Cohesion	▪ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[First Step] Modelling the Domain

- Defines fundamental entities, constructs, and their relationships
- Leverage existing domain knowledge
- External domain knowledge – we are not the first...
 - Enterprise Integration Design Patterns
 - [Ref. Enterprise Integration Patterns, Gregor Hohpe & Bobby Woolf, 2003]
 - Enterprise Application Design Patterns
 - [Ref. Patterns of Enterprise Application Architecture, Martin Fowler]

Translating the Architecture

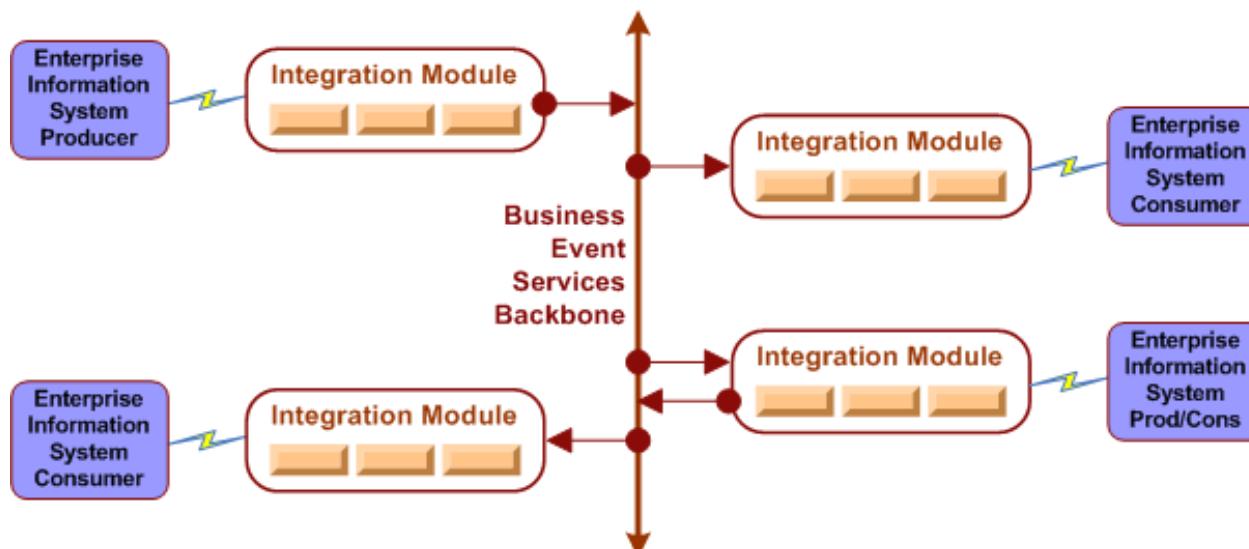
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
▪ Separation of Concerns	▪ Tight Cohesion	▪ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Integration Modules

- Provides a logical grouping of business operations as a single integration point
- Provide either a source, target, or bi-directional business flow.



Translating the Architecture

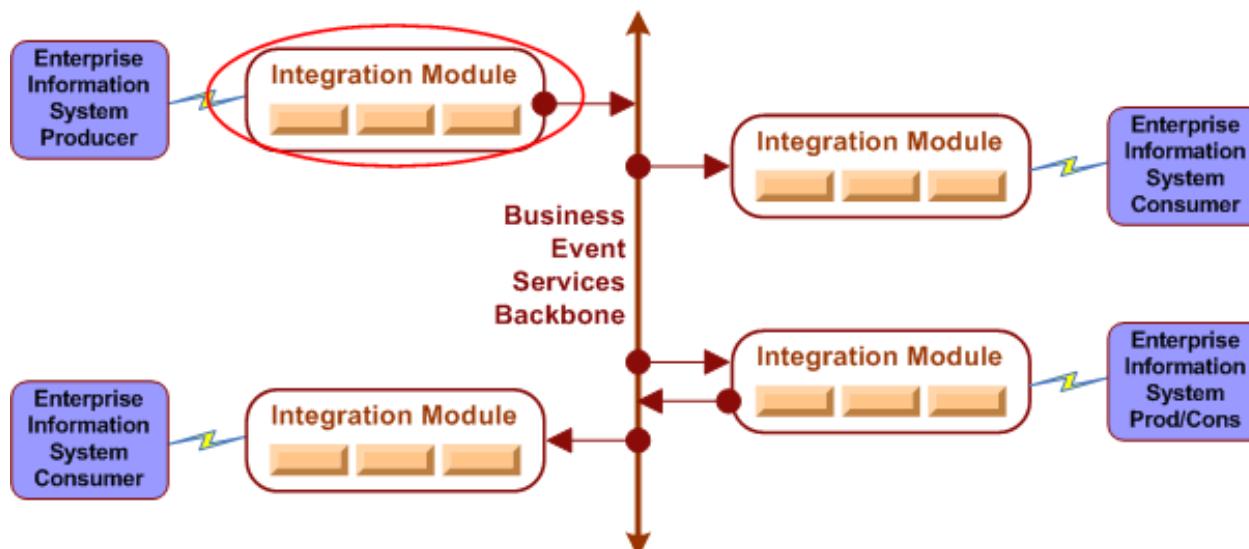
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
▪ Separation of Concerns	▪ Tight Cohesion	▪ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Integration Modules

- Provides a logical grouping of business operations as a single integration point
- Provide either a source, target, or bi-directional business flow.



Translating the Architecture



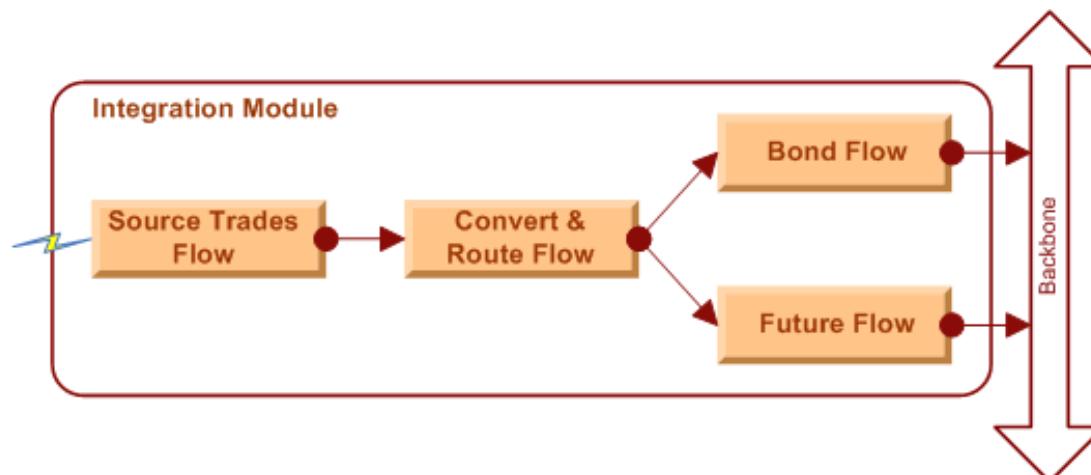
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Flows

- Provide cohesive operations on a business artifact as a synchronous operation
- Multiple flows can be chained to isolate concerns
- Standard event container allows any data type to be transported



Translating the Architecture



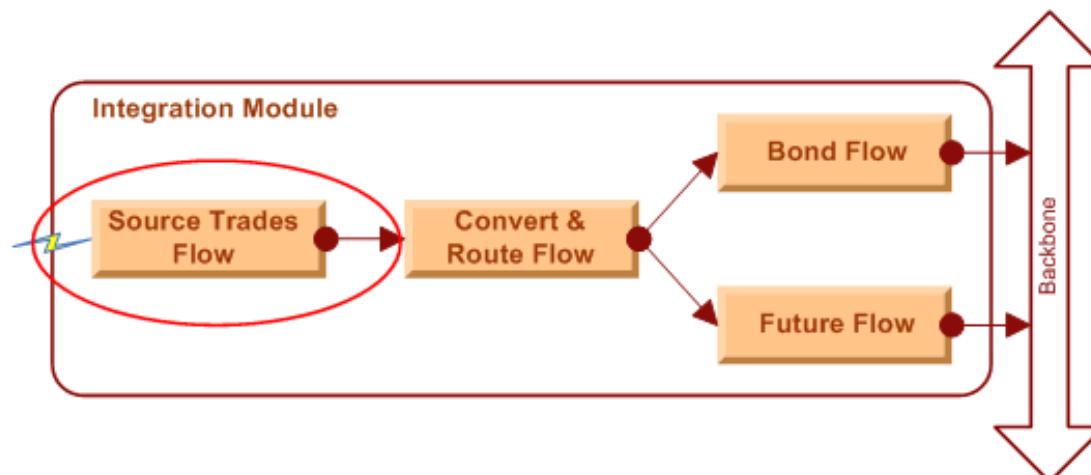
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
▪ Interchangeable Constructs	▪ Highly Testable Constructs	▪ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Flows

- Provide cohesive operations on a business artifact as a synchronous operation
- Multiple flows can be chained to isolate concerns
- Standard event container allows any data type to be transported



Translating the Architecture



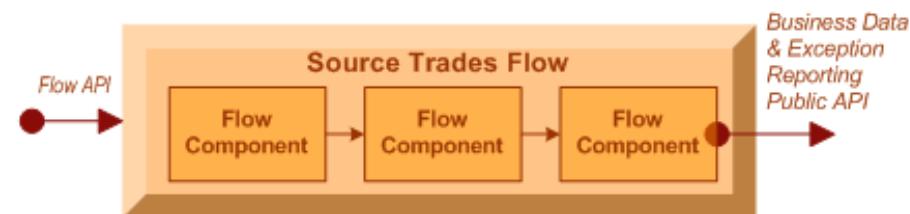
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
▪ Control & Management	▪ Monitor & Alerting	▪ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Flow Components

- POJO implementations provide independent, interchangeable operations
- Isolated component white box testing and chained black box testing
- All flow components within an atomic operation guaranteeing data integrity



Translating the Architecture



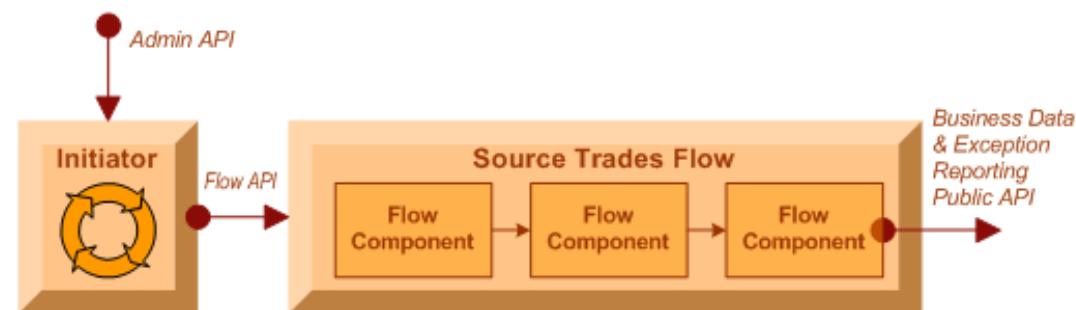
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
▪ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Initiators

- Business flow control (start & stop)
- Notification of flow status (running, stopped, recovery, stoppedInError)
- Configurable exception handling for automated flow recovery
- Public API for external client management



Translating the Architecture

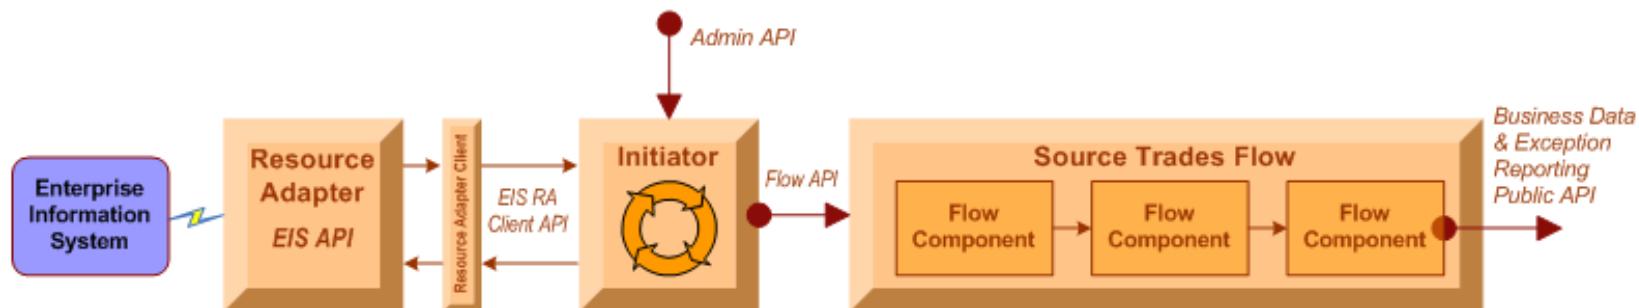
[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
✓ Support any Protocol/Entity	▪ Business Event Tracking	▪ Operational Audit

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Resource Adapters

- Support any type of protocol / API
- Extend guaranteed integrity to non-guaranteed EIS sources/targets
- No proprietary knowledge required; only EIS API
- Resource adapters pluggable re-usable components (SUN JCA spec)
- 3rd party off-the-shelf or internally developed



Translating the Architecture



[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
✓ Support any Protocol/Entity	✓ Business Event Tracking	✓ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Administration Console

- Initiator management
 - runtime status (stopped, running, recovery, stoppedInError)
 - runtime control (start, stop, start-up type {auto, manual, disabled})
- Flow management
 - flow component chain view and drill down
 - audit log of flow actions (stop/start)
- Wiretap management
 - dynamic creation/deletion of real-time in-flight data capture
- User management
 - creation, deletion, and permission maintenance

Translating the Architecture



[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
✓ Support any Protocol/Entity	✓ Business Event Tracking	✓ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] User Console

- In-flight business event searches
 - by module(s)
 - by flow(s)
 - by business stream(s)
 - parameter filters
 - business content, date/time, names, etc
- Event “Life Identifier” linking
 - flow component chain view and drill down
 - audit log of flow actions (stop/start)
- User management
 - creation, deletion, and permission maintenance

Translating the Architecture



[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
✓ Support any Protocol/Entity	✓ Business Event Tracking	✓ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] Monitor Clients

- Big Brother / Hobbit client
- Email client
- Log scrapper
- HTTP client

Translating the Architecture



[Key Strategies]

✓ Single Point of Integration	✓ Loose Coupling	✓ Variety of Standard Contracts
✓ Separation of Concerns	✓ Tight Cohesion	✓ Standard Design Patterns
✓ Interchangeable Constructs	✓ Highly Testable Constructs	✓ Data Integrity
✓ Control & Management	✓ Monitor & Alerting	✓ Minimal Manual Intervention
✓ Support any Protocol/Entity	✓ Business Event Tracking	✓ Operational Audit

- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

[Construct] OSS Test Frameworks & Ikasan FlowTestHarness

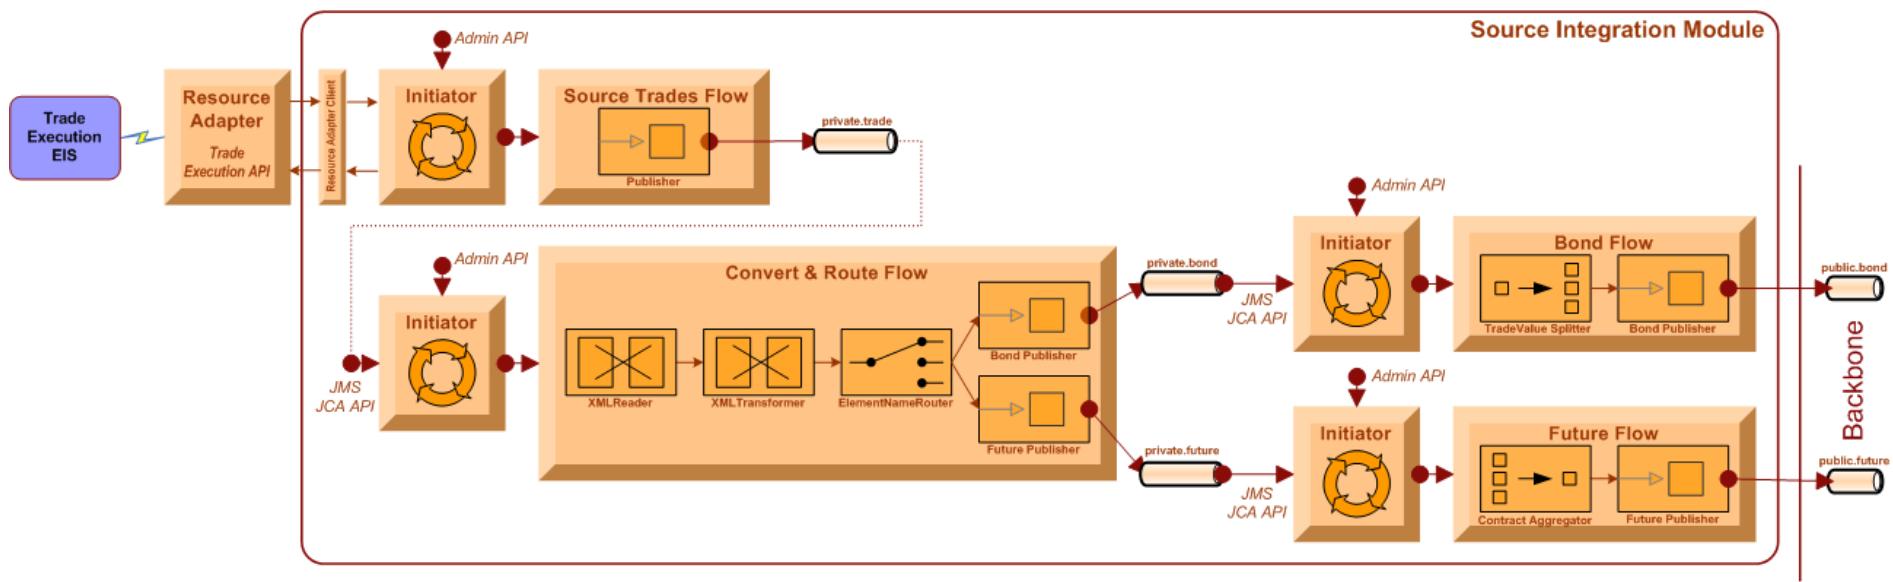
- Flow component automated unit testing
 - Junit, JMock, XMLUnit
- Initiator automated unit testing
 - JUnit, JMock
- Flow automated unit (integration) testing
 - Spring-test, Ikasan FlowTestHarness, HSQL in memory database
- Module integration testing
 - Currently manual
 - Future: automation through maven plugins

Translating the Architecture



- Sample Source Integration Module
 - Source Trades Flow
 - Get trades (CSV format) from EIS with minimal failure potential
 - Convert & Route Flow
 - Transform CSV -> XML syntax; Transform XML semantics; Route on type
 - Bond Flow
 - Split one bond event into many bond events based on trade size rule
 - Future Flow
 - Aggregate multiple future contract events until x then push as single event

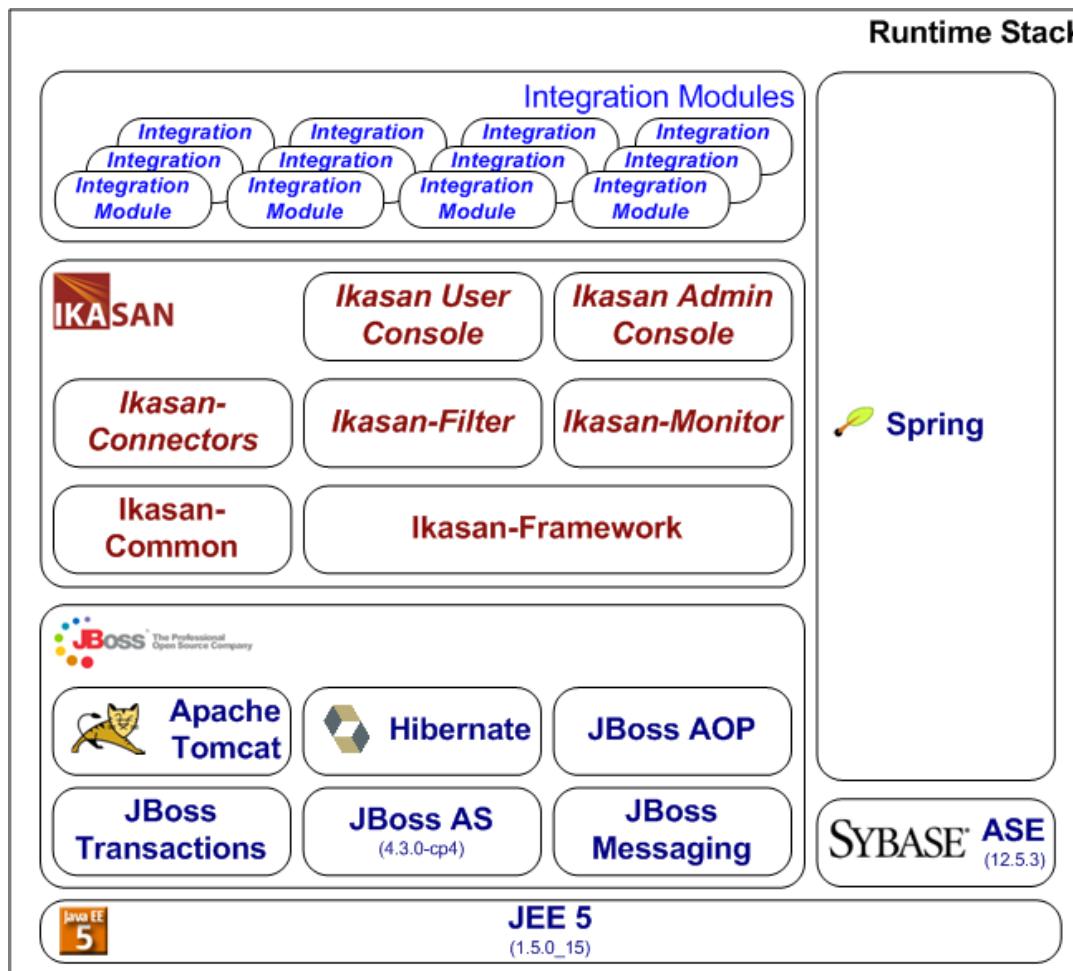
- Problem Domain
- Architecture Strategies
- **Translating Architecture**
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary



Technology Stack



■ Runtime Stack



- Problem Domain
- Architecture Strategies
- Translating Architecture
- **Technology Stack**
- Features
- Development Lifecycle
- Roadmap
- Summary

Technology Stack



- Development Stack
 - Core technologies
 - Java 1.5+ / JEE
 - XML
 - Spring Framework 2.6
 - Test Frameworks
 - JUnit
 - JMMock
 - SpringTest
 - Ikasan Flow Test Harness 0.7.10
 - Scaffolding
 - Source Management – Subversion
 - Requirements and Issue Management – JIRA
 - Peer Review and Acceptance Management – Crucible, FishEye
 - Project Build Management – Maven
 - Project Artefact Management – Sonatype Nexus
 - Continuous Integration – Hudson

- Problem Domain
- Architecture Strategies
- Translating Architecture
- **Technology Stack**
- Features
- Development Lifecycle
- Roadmap
- Summary

Features



- Version 0.7.10 current production release
- Interoperability
 - Protocols
 - JMS
 - FTP (transactional & chunking support)
 - SFTP (transactional & chunking support)
 - RDBMS (via Hibernate ORM)
 - SMTP (Email)
 - HTTP(s)
 - Vendor APIs
 - Trax Messaging
 - PATS Trade Feed
 - Espeed Trade feed
 - Tradeweb Trade Feed
 - Bloomberg Data License Per Security
 - Data Constructs
 - Event – Ikasan internal entity container
 - Payload – actual entity in-flight
- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Features



- **Functionality**
 - Initiators – business flow control
 - Allows a client to stop, start, enable, disable, query runtime status
 - Event Driven Type
 - Any resource adapter protocol
 - Schedule Driven Type
 - Quartz Scheduler for absolute and relative invocation scheduling
- **Flows**
 - Full Transactional Semantics
 - Non-transactional
 - Local transactions
 - 2-Phase transactions
 - Last resource commit optimisation (LRCO)
 - incorporate non-transactional resources within a distributed transaction
 - Configurable Automated Recovery
 - Rollback operations & stop
 - Rollback operations and retry in x seconds
 - Rollback operations and retry a number of times or indefinitely
 - Report the exception and continue
- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- **Features**
- Development Lifecycle
- Roadmap
- Summary

Features



- **Functionality**
 - **Flow Components**
 - Transformers
 - XSLT
 - XML Reader
 - XML Validator
 - any POJO implementing Transformer interface
 - Routers
 - Recipient list
 - Content
 - Xpath
 - Any POJO implementing Ikasan router interface
 - Sequencers
 - Aggregator
 - Splitter
 - Any POJO implementing Ikasan sequencer interface
 - Endpoints
 - Loggers
 - Hard wiretap
 - Any POJO implementing Ikasan endpoint interface
 - Any resource adapter client API
 - Problem Domain
 - Architecture Strategies
 - Translating Architecture
 - Technology Stack
 - **Features**
 - Development Lifecycle
 - Roadmap
 - Summary

Features



- **Functionality**
 - Dynamically capture runtime data at any point in the flow
 - Users able to search on business content across a flow
- **Problem Domain**
- **Architecture Strategies**
- **Translating Architecture**
- **Technology Stack**
- **Features**
- **Development Lifecycle**
- **Roadmap**
- **Summary**

Development Lifecycle



- Approach based on TDD/BDD Agile dogma
 - Gather/refine user stories and create use cases
 - Sketch the stories into a module/flow design
 - Create/refine the Flow Test Harness reflecting the design
 - Create modules/flows/components to meet test harness acceptance
 - Peer/pair walk through (10 minute daily stroll)
 - Design reviews
 - Code reviews
 - Continuous Integration (CI) test coverage increases
 - Deliver functionality to users for proofs
 - Repeat!
- Recent Example Implementation Times
 - 3 days – Completed module delivery for transformed XML; real-time and batch aggregation; delivered over HTTP(S) to downstream application
- Demo

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- **Development Lifecycle**
- Roadmap
- Summary

Roadmap



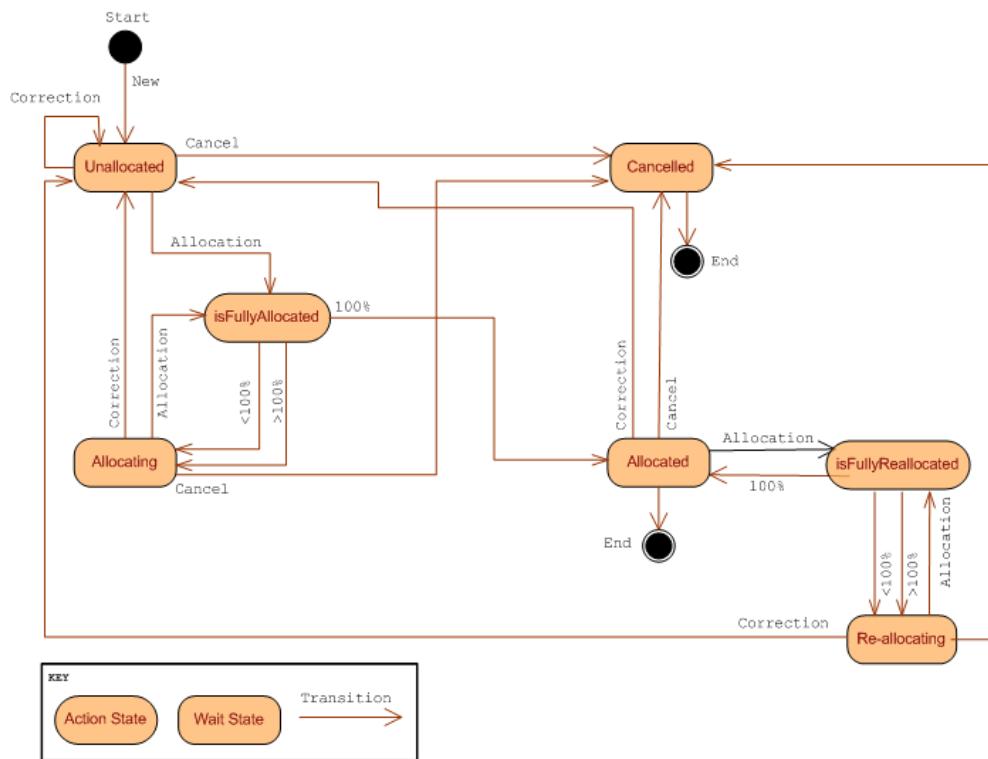
- Version 0.8.0 due shortly, currently in beta
- New Features
 - Event hospital
 - Resubmission of failed business events
 - Public API invocable over HTTP
 - Full history of exceptions on a given event
 - Simplified Domain Model
 - Support any user defined entity in the domain model
 - Support of any user defined domain properties

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Roadmap

- Version 0.8.x or 0.9.0
 - New Feature
 - State Modeling
 - Standard Event Condition Action (ECA) State Pattern supporting
 - wait states
 - action states
 - conditional & default transitions

- Problem Domain
 - Architecture Strategies
 - Translating Architecture
 - Technology Stack
 - Features
 - Development Lifecycle
 - **Roadmap**
 - Summary



Roadmap



- Future Releases
 - Functionality
 - event replay
 - business subscription streams as an exposed registration service
 - Monitoring & Management
 - Flow status notification management service
 - metrics and stats
 - active monitoring rules
 - Documentation
 - Quick Start Guide
 - Developer Guide
- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary

Summary



- Ikasan Project Resource Summary
 - Web Site: <http://www.ikasan.org>
 - SVN: <http://open.jira.com/svn/IKASAN>
 - Wiki: http://sourceforge.net/apps/mediawiki/ikasaneip/index.php?title=Main_Page
 - JIRA: <http://open.jira.com/browse/IKASAN>
 - Mailing List: <http://lists.sourceforge.net/lists/listinfo/ikasaneip-user>
 - IRC: <http://sourceforge.net/apps/mediawiki/ikasaneip/index.php?title=IRC>

- Problem Domain
- Architecture Strategies
- Translating Architecture
- Technology Stack
- Features
- Development Lifecycle
- Roadmap
- Summary



Open Source Enterprise Integration Platform

*Jeff Mitchell
ikasan.org
July 2010*